

KUNAL SACHDEV

437-663-6032 Email Portfolio LinkedIn GitHub Credly

EDUCATION

University of Waterloo

Bachelor of Computer Science, Minors in Statistics and Economics

Sep. 2023 – Present

Waterloo, ON

EXPERIENCE

Full Stack Software Developer - Cloud and ML Co-op

Nokia Canada

January - April 2026

Kanata, ON

- Incoming - Cloud and Network Services (CNS) team. AI/ML solutions for global networks.

Data Scientist Co-op

Ontario Lottery and Gaming Corporation

Sept 2025 – Present

Toronto, ON

- Deployed a **multimodal RAG** chatbot using **LangChain**, **Azure AI Search**, and **OpenAI API (GPT-4o)** to efficiently query insights from business documents, saving business teams **215+** hours per month.
- Fine-tuned an **XGBoost** classifier using **Databricks**, **Hadoop**, and **PySpark** to segment **115M+** players into 6 categories, enabling targeted ad campaigns. Collaborated with stakeholders to design **A/B tests**, optimizing campaign engagement.
- Worked in an **Agile** environment, following **SDLC** best practices, creating **PRs**, and automating deployments using **CI/CD** pipelines (Jira, **GitHub**, TeamCity, **Docker**, Octopus Deploy).

GenAI Engineer Intern

Edelweiss Life Insurance

May 2025 – August 2025

Mumbai, India

GenAI-powered Business Insights with Visualization

- Built a serverless **RAG** app on **AWS Bedrock** enabling real-time conversational access to AWS Data Lake, replacing **200+** static BI dashboards and enabling **Rs. 300,000+** in estimated annual operational cost savings.
- Fine-tuned **CodeLlama-7B-Instruct** from **HuggingFace** using QLoRA on NVIDIA A10 GPUs; deployed model to AWS **SageMaker** with **TensorRT-LLM**, accelerating inference by **1.8x** and cutting memory usage by **50%**.
- Engineered **Lambda** functions for query validation, **Redshift** schema extraction, OpenSearch retrieval, and LLM inference; achieved **99%** API uptime and reduced **8,700+** annual emails per user via automation.
- Boosted SQL generation accuracy by **28%** using **Cohere** embeddings + **OpenSearch** retrieval for **few-shot prompting**. Integrated **API Gateway** + **Streamlit** UI, serving **1.5K+** daily active users with real-time tables/visuals.

Data Science – Policy Grievance & Mis-selling Prediction

- Productionized a **LightGBM** + **CATBoost** ensemble to flag grievance/mis-selling risk across **809K+** policies, surfacing **96.9%** of grievance cases in the top decile, empowering agents to preempt complaints.
- Developed a scalable ML pipeline on AWS **SageMaker**, containerized with **Docker** on AWS **ECR**, and automated weekly predictions with **EventBridge** Scheduler cron job, reducing prediction latency by **66.7%**.

Data Science Intern

Info Origin Inc.

May 2024 – August 2024

Remote

- Deployed a custom **spaCy NER** model with Streamlit UI to auto-extract key resume entities, cutting HR candidate screening time by **40%** and streamlining early-stage hiring workflows.
- Optimized hyperparameters for a custom **5-layer** neural network using Google's NNLM embeddings and **PyTorch TensorDataset/Dataloader** for efficient batching/shuffling of **2,200+** articles, achieving **96.4%** accuracy in news classification

PROJECTS

WLP4 Compiler | C++, Bash, Linux, GNU, GCC, GDB

- Built a compiler for WLP4, a subset of C, implementing Simplified Maximal Munch for **scanning**, LR(1) for **parsing**, context-sensitive analysis, and **MIPS** assembly generation, with **30%** improvement in performance using optimization techniques.

Car Dealership App | Architecture, Node.js, Express, MongoDB, Flask, Django, React, Kubernetes | GitHub

- Architected a full-stack microservices app with a **React** frontend, and **3** services: **Django** + **SQLite** backend, reviews REST APIs with **Express.js** + **MongoDB**, and a sentiment analyzer on IBM Cloud; CI/CD with GitHub, & deployment on **Kubernetes**.

Biquadrus | UML Diagram, C++, XWindows, Git, Cygwin, Bash, Linux, GNU, GCC, GDB | GitHub

- Implemented a two-player Tetris-style game in **C++** using modular **OOP**, applying **Factory** and **Observer** design patterns to support **6+** blocks, **3+** levels, **3** special actions, and real-time text/graphical rendering via XWindows.

Winning the Space Race with Data Science | BeautifulSoup, Db2, Matplotlib, Plotly, Scikit-learn | GitHub

- Developed an ML pipeline predicting Falcon 9 landing success with **87.5%** accuracy, using GridSearchCV-tuned **Decision Tree**; automated EDA of **90+** launches via **Db2**, SpaceX API, **BeautifulSoup**, and designed **Plotly/Folium** dashboards.

TECHNICAL SKILLS

Languages/Cloud: AWS, IBM Cloud, Python, C, C++, R, SQL, JavaScript, HTML/CSS, Tailwind CSS

Machine Learning: Hugging Face, Pandas, NumPy, Scikit-learn, PyTorch, TensorFlow, Transformers, Jupyter, BeautifulSoup

Web Dev/Databases: React, Next.js, Express, Node.js, Flask, Django, Dash, SQLAlchemy, MongoDB, MySQL, SQLite, Db2, Oracle

Technologies: Git, Docker, Kubernetes, Linux, bash, GNU, GCC, GDB